

CS44800 Project 3

Relational Operators, Joins and Query Evaluation Pipelines

Spring 2020

- *Unlike project 2, this project is an INDIVIDUAL project and not a group project.*
 - Total Points: 100
 - Due: 11:59PM EST, Friday, April 10, 2020. Submit using Blackboard.
 - (There will be a 10% penalty for each late day. After 5 late days, the project will not be accepted.)
-

Learning Objectives

1. Understand how Relational Algebra operators are implemented inside of a database system
2. Understand algorithms for executing different Relational Algebra operators
3. Implementing and understanding join operations from a lower level perspective

Background

The goal of this project is to implement Relational Algebra operators using the Iterator interface inside a working database system. Like Project 2, this project is based on Minibase, a small relational DBMS, structured in several layers in order to allow a modular, abstract approach to implementing a DBMS

As you have learned in class, a typical database query processor optimizer breaks down queries into trees of relational operators, implemented as iterators. The iterators (i.e., [HeapScan](#) and [HashScan](#)) deal with file access directly, and return records or their ids. In this project, you will build and use a higher-level view of these records with the provided classes [Schema](#) and [Tuple](#).

When executing an SQL query, a DBMS transforms the query into an equivalent Relational Algebra expression. Each Relational Algebra operator is implemented through an *iterator interface*. This interface consists of the following functions:

- open()
- close()
- hasNext()
- getNext()

Relational Algebra operators are then chained together to form a Query Tree. A query is then evaluated by starting with the root operator of the tree and calling **hasNext()** and **getNext()** to recursively request the next tuple from its child Iterators until all rows have been exhausted.

This results in an implementation where the results of a query are evaluated lazily and in a non-blocking fashion. Any given Iterator only requests results from its child Iterators when necessary and in most cases does not need to compute its entire set of results before generating results for the Iterators higher up the query tree – tuples simply “flow” up the query tree as they are evaluated. This *pipelined* approach results in queries that can start generating results immediately, and do not require any storage space for intermediate results.

In this project, you will be implementing Relational Algebra Operators using this Iterator interface, creating the functionality necessary to actually execute any arbitrary query in Minibase.

You can find the API for the Minibase classes you will be using [here](#). Please familiarize yourself with the API before starting this project. In particular, the classes you will need to use are:

- Classes in the *relop* package,
- HeapFile and HeapScan in the *heap* package,
- HashIndex, HashScan and BucketScan in the *index* package, and
- AttrOperator and AttrType in the *global* package.

Part 1: Scan Operators - From Records to Tuples [30 points]

All Relational Algebra operators you will be implementing in this project will be subclasses of the Iterator abstract class. The Iterator class provides a *schema* field (of protected visibility) that can store the schema of any output tuple from that Iterator, as well as defines the following methods that must be implemented:

- public void restart()
- public boolean isOpen()
- public void close()
- public boolean hasNext()
- public Tuple getNext()

The constructor for each Iterator is equivalent to the isOpen() function of the Iterator interface.

Note: you do not have to implement the explain() method

Your first task is to construct the initial scan operators. These are Iterators that read through an entire relation or index and report one Tuple at a time. Minibase already has methods defined to scan through and retrieve records in raw byte format from the HeapFile and HashIndex classes, but we need to create Iterator versions of these scans that will act as wrappers for these scans and will return Tuples with defined schemas as opposed to the raw byte information.

Specifically, we require the following Scan operators:

- **FileScan**
A *HeapScan* that returns *Tuples* instead of byte[] 's

- **KeyScan**
A *HashScan* that returns *Tuples* instead of *RIDs*
- **IndexScan**
A *BucketScan* that returns *Tuples* instead of *RIDs*

The FileScan implementation is already provided to you. You will need to implement similar scan operators for the KeyScan and IndexScan. Some useful hints and tips:

- Each constructor should initialize the inherited field *schema*. (i.e. it is given as a parameter to these three iterators)
- Don't be surprised by how little code these classes require
- *HashScan* scans the hash index for records having a given search key. *BucketScan* scans the whole hash index. Those classes only return RIDs. You have to build wrapper classes KeyScan and IndexScan that return Tuple.

Part 2: Primitive Operators [30 points]

Now that you have the basic leaf nodes of most query trees, you can make some more interesting iterators. Your next task is to implement the three fundamental operations of relational algebra:

- **Selection**
Filters the output of another Iterator on a set of Predicates. When there are multiple Predicates, they are assumed to all be Ored together (AND is implemented through chaining separate Selection operators together).
- **Projection**
Performs a projection on a Tuple by removing columns from the output of another Iterator. This will change the Schema of all Tuples output by the Projection. Note: you should not remove duplicates.

Part 3: HashJoin [40 points]

For the last part of the project, you are required to implement a **Hash-Join**. Please read the relevant chapter from your textbook carefully before proceeding with this part. Furthermore, the source code for the "Nested Loops Join" is provided for you in SimpleJoin.java. You may consider studying the source code and use it as a guidance. For example, this join was implemented using a pre-fetching approach, which we recommend you follow for the primitive operators you have to implement.

A Hash-Join consists of two stages:

- First, partition phase, each input relation is partitioned in k partitions or buckets using the same hash function. Luckily, you do not have to implement most of this from scratch – *you can make use of the IndexScan within Minibase to perform this.*

- Then, probing phase, for each bucket on the smaller input relation, we build an in-memory hash table with its tuples. Next, this hash table is probed using the tuples from same bucket in the other relations. This process is repeat with all the buckets.

For the first phase of a Hash-Join, you need to transform the input Iterators into an IndexScan with the join column as the key. **Note: You cannot simply cast the child Iterators of the HashJoin as an IndexScan. You must construct a new temporary HashIndex and HeapFile populated with the results of the child Iterator and create an IndexScan.**

Once the child Iterators have been transformed into IndexScans with the join column as the search key, you can iterate over the IndexScans to determine matching results. An IndexScan iterates over an index **bucket-by-bucket**. Note that the contents of a bucket may be in any arbitrary order and may not all be matches, even if all entries have hashed to the same bucket. Therefore, you will have to track matching tuples within a bucket as you iterate through. To facilitate this, you are provided with a class **HashTableDup**. This class acts as a HashTable, but allows for multiple identical keys to exist. A common pitfall is to not comprehend that the matching is done bucket by bucket. Therefore, a tip is to regenerate the hash table when the program has completed with searching one bucket. The rationality behind this approach is that

The benefit(s) and reasoning for using Hash Join is briefly mentioned right after the pseudocode in the next section.

Hash-Join Related Notes

(You are required to read up on the actual algorithm from your textbook)

We want to join tables R and S on their common p table/column as follows:

$$R \bowtie_p S$$

Let's assume we want to join tables R and S with a equijoin predicate p .

Some guidelines on implementing Hash Join are provided below:

1. As mentioned earlier in the handout, your first task would be to transform both iterators to IndexScans to facilitate scanning through the records. By using the IndexScan iterator, all the records from the relation will be partitioned in buckets by applying a hash function h to each tuple. This first step is called **partition stage**.
2. Ideally you would keep the table with smaller cardinality (let's assume R is the smaller relation) hashed as key value pairs in a hash table. The other table (i.e. S in this case) would also be probed against this hash table.
3. The second phase is the **probing stage**, where you will match the tuples for the relations R and S , bucket by bucket. Then, you will insert the tuples from bucket i into the in-memory hash table and probe the tuples of the S of the same bucket (i). The rationality behind is that the memory will not be enough to fit all the tuples from S ; instead, we are only keeping in memory the tuples of the bucket we are currently probing.
4. Whenever you will find a match, store the tuple and continue.

The basic advantage of Hash Join is the fact that we only need one full scan of one of the tables/relations (in our example – the S relation) to find the output. This is almost double as efficient as nested loops join (except it is only limited to equijoins). The same join can be done with nested loops join but the memory overhead and the time overhead – both would be double. Therefore, we emphasize on your correct implementation of the Hash Join for acquiring full grade.

For the technical capabilities of the hash functions and hash code comparisons, we strongly encourage you to study the *IndexScan*, *SearchKey*, *Tuple* (particularly the join method), *Schema*, *HashJoin* and *HashTableDup* classes in the Minibase javadocs. You will need all of them for the complete implementation of the project. Once you understand the structure and the requirements, the coding is quite simple. But understanding your goal and the classes at your disposal is the most important part going forward.

PART 4: Query Evaluation Pipelines [UNGRADED]

Now that you have a set of operators, you are ready to use these operators to form query evaluation pipelines (QEPs, for short) that can evaluate some queries. *This is important for you to see what you implemented in the previous parts of the project in action.*

To simplify your task, you are given the code for some simple QEPs that correspond to some simple queries. You will find that code in ROTest.java. You have to study that code in order to understand how a schema can be created, data can be inserted, and operators can be connected in Minibase.

Consider the relational schema of two tables:

- **Employee (EmpId, Name, Age, Salary, DeptID), and**
- **Department (DeptId, Name, MinSalary, MaxSalary)**

For simplicity, assume one-to-one relationship between the two tables.

The first four are given to you for demonstration purposes. You are encouraged to implement the queries 5 through 8.

You are required to programmatically form a QEP for each of the queries given below:

1. Display for each employee his Name and Salary
2. Display the Name for the departments with MinSalary = 1000
3. Display the Name for the departments with MinSalary = MaxSalary
4. Display the Name for employees whose Age > 30 and Salary < 1000
5. For each employee, display his Salary and the Name of his department
6. Display the Name and Salary for employees who work in the department that has DeptId = 3

7. Display the Salary for each employee who works in a department that has MaxSalary > 100000
8. Display the Name for each employee whose Salary is less than the MinSalary of his department

Note that whenever you need to connect a join operator to your QEP, you can use either the hash-join operator you implemented, or the nested-loops join operator that is already given to you.

Following the paradigm in ROTest.java, you are encouraged to write **one test for each of the above queries** (queries 1-4 are worked out for you to look up) inside QEPTTest.java.

Hints

Some useful hints and tips:

- For Relational Algebra operators that may not always find results (i.e. Selection and HashJoin), you need to be careful about exactly how you are determining the result of the hasNext() method call. Unlike with the basic scans, you cannot simply just return the result of the child iterator's hasNext() method. You may check the SimpleJoin.java class for an example of how to properly handle this.
- The join column specification for HashJoin is based on the schemas of the individual child Iterators *before* the join. For example, a HashJoin constructor call with “(...,0,2)” would mean the join columns are column 0 of the left Iterator and column 2 of the right Iterator.
- Implementing the HashJoin will likely be the most complicated part of the project. Please budget your time accordingly.
- Please keep an eye out on Project 3 Notes periodically – as the topmost pinned thread on Piazza for any update or announcement regarding the project.

Running the Project

For the convenience of testing for you, this project is made on top of Maven framework.

Building with Maven

We have made Maven available on the **university machines**. In order to use Maven, you will have to set the Path variable through the terminal as follows (*you only need to do this once*):

```
export PATH=/homes/afiroze/cs448/apache-maven-3.6.0/bin/:$PATH
```

To set up Maven in **your own machine**, please follow the instructions [here](#). In the terminal, navigate to your application directory (the one with the pom.xml file). Input the following command:

```
mvn clean compile assembly:single
```

Maven should now build your project. After the build finishes, you will find the compiled class files in the `target` directory.

Upon making the binaries, you can use the following command to run all the test cases on your implementation.

```
mvn test
```

Remember to run this command from the parent project directory.

Moreover, the testcases are also provided as a JUnit suite inside the test directory in **src/test/java/tests**. The testcases consist of a mix of relational operators combined through pipelines. These testcases can be used to verify the correctness of your implementation for each of the parts described in this handout. In the specific case of the **Hash-Join**, you must follow the guidelines in this handout – *iterating through the HashIndexes bucket by bucket* – to be awarded the full credits. To validate this, we will manually inspect your submission.

The input relations, used by the testcases, are defined on a pair of files for each table, a JSON with the schema and CSV with the corresponding data. Likewise, a text file, `solution.txt`, includes the expected solution for each testcase. You can play around with the input data located in **/src/test/resources** and accordingly modify the resultant records in the **solution.txt** file in your parent project directory. Make sure that you update in both places correctly if you are adding in your own data. This is quite helpful for a sanity check of testing your system with queries.

Note: The structure of solution.txt and the data in part 4 are not in the same format. Solution.txt represents each row/tuple as a string delimited by “|” (pipe) and fields by “:” (colon). On the other hand, Employee.txt and Department.txt has their rows/tuples delimited by a newline separator and fields delimited by comma.

This project has been tested and run through the terminal on the university Linux machines. We cannot provide support for using different IDEs or running the program on different platforms, but some posts have been made on Piazza regarding setting up IDEs to run the Minibase projects that you can reference.

Skeleton Code and Submission Instructions

Skeleton of the code has been made available to you in the same zip file as this handout, and the documentation is provided in the `javadocs` directory.

Note that this code skeleton is a complete starting point for the project, i.e., the *bufmgr*, *heap*, and *index* packages are provided for you (in jar files).

You should turn in your code and a `readme` file. All files need to be zipped in a file named: **your_career_login_ro.zip**.

In the readme file, put anything you would like us to know. We should be able to compile/run your program using the Maven available at the CS department Unix machines.

Do not change the directory structure of the code. The directory structure of your zip file should be identical to the directory structure of the provided zip file. Your grade may be deduced 5% off if you don't follow this.

These files should be submitted on Blackboard.